

AGENTIC INFRASTRUCTURE • FIELD NOTES

Managing long-running agents

Leave an agent alone for 300 hours and you'll end up auditing its coping mechanisms instead of its engineering.


SWIPE →
@basho-parks

islandmountain.io

THE MECHANISM

Adaptation accumulates inertia.

Agents adapt continuously. That's the point of them. But every workaround, retry policy, validation check, and defensive heuristic gets folded into the agent's behavioral memory.

A decorative horizontal line consisting of 25 vertical bars of varying heights, colored in a light orange or tan hue, spanning the width of the page.

THE DRIFT

**It becomes the
sum of every
problem it ever **hit.****

Given enough time, the agent stops behaving like the engineer you assigned to the task and starts behaving like its own operational history. The codebase turns into a record of survival.

THE TRAP

This failure doesn't **look** like failure.

Tests still pass. Pull requests still open.
Nothing obvious flags that the system is
now optimizing for surviving its own
history instead of solving the problem
you directed it to solve.

FIX 01

of 03

Keep an immutable Hour Zero.

Freeze a behavioral baseline the moment the session begins. Never overwrite it. Compare today's agent only to yesterday's and you're measuring decay with a ruler that changes length every time you pick it up.

FIX 02

of 03

Force the agent to leave receipts.

Human engineers leave commit messages and context. Hold agents to the same standard. If a retry strategy changes because an API threw repeated 429s, that trigger gets logged beside the behavioral change.

Fingerprint the **behavior**.

Ops already watches CPU, memory, and latency. Those are machine properties. Autonomous systems add agentic properties. None proves the agent is wrong. Together they show it's becoming different.

- Verbosity

- Retry density

- Error tone

- Decision latency

THE DOCTRINE

Drift is not preventable. It's the cost of **adaptation.**

The job is making drift visible, explainable, and reversible. The controls that hold live outside the model: frozen baselines, logged triggers, tracked fingerprints. Enforced by infrastructure, not requested by a prompt.

FIELD NOTES • WEEKLY

I write about agentic
infrastructure
and governance
every week.

Follow for the next one.